

XGBoost Model Comparison Using Python and R

Introduction

This report presents a comparative analysis of XGBoost implementations across Python (via scikit-learn) and R (using both direct xgboost() and caret packages) using datasets of varying sizes. The goal is to assess and compare predictive performance and computational efficiency of each approach on synthetic data generated through sampling with replacement. Predictions are made using logistic regression models.

Methodology

Tools Used: Python (scikit-learn), R (base xgboost, caret), 5-fold cross-validation.

Dataset: Synthetic data sampled from original dataset; predictions made using logistic regression.

Evaluation Metrics:

- Accuracy on the testing set.
- Time taken for model fitting in seconds.

Results Table

| Method | Dataset Size | Accuracy | Time (s) |
|-----------------------|--------------|----------|----------|
| Python XGBoost | 100 | 0.8500 | 0.88 |
| Python XGBoost | 1,000 | 0.9380 | 1.88 |
| Python XGBoost | 10,000 | 0.9737 | 7.99 |
| Python XGBoost | 100,000 | 0.9885 | 4.02 |
| Python XGBoost | 1,000,000 | 0.9918 | 10.46 |
| Python XGBoost | 10,000,000 | 0.9932 | 94.28 |
| R xgboost() direct CV | 100 | 0.9009 | 0.34 |
| R xgboost() direct CV | 1,000 | 0.9340 | 0.65 |
| R xgboost() direct CV | 10,000 | 0.9592 | 1.44 |

XGBoost Model Comparison Using Python and R

| | | | |
|-----------------------|------------|--------|---------|
| R xgboost() direct CV | 100,000 | 0.9720 | 7.96 |
| R xgboost() direct CV | 1,000,000 | 0.9749 | 72.60 |
| R xgboost() direct CV | 10,000,000 | 0.9756 | 238.84 |
| R caret xgboost | 100 | 0.9100 | 1.88 |
| R caret xgboost | 1,000 | 0.9349 | 0.70 |
| R caret xgboost | 10,000 | 0.9724 | 1.64 |
| R caret xgboost | 100,000 | 0.9847 | 12.97 |
| R caret xgboost | 1,000,000 | 0.9887 | 109.28 |
| R caret xgboost | 10,000,000 | 0.9896 | 1007.01 |

Analysis & Recommendation

While all methods show improvement in accuracy with larger datasets, Python's implementation of XGBoost provides consistently high accuracy with better computational efficiency at scale. For instance, Python XGBoost reaches 0.9932 accuracy at 10 million rows in 94 seconds, whereas R caret takes over 1000 seconds for slightly lower accuracy. Therefore, Python XGBoost is recommended for large-scale applications due to its performance and speed.

Conclusion

XGBoost remains a powerful model across platforms. However, Python's implementation is better optimized for speed and large data. R caret provides more interpretability and structured CV, which is ideal for smaller projects. Selection depends on use case and resource constraints.